PCG for the infinite-dimensional mode-k subproblem (missing data)

We consider the linear system in the unknown $W \in \mathbb{R}^{n \times r}$

$$\left[(Z \otimes K)^\top SS^\top (Z \otimes K) + \lambda(I_r \otimes K) \right] \mathrm{vec}(W) = (I_r \otimes K)\, \mathrm{vec}(B), \tag{1}$$

where $K \in \mathbb{R}^{n \times n}$ is a (symmetric) psd kernel matrix, $Z \in \mathbb{R}^{M \times r}$ is the Khatri–Rao product of the other factors, $S \in \mathbb{R}^{N \times q}$ selects the q observed entries (so $S^\top \mathrm{vec}(T)$ equals the observed values), and $B = TZ \in \mathbb{R}^{n \times r}$. Throughout we assume $n, r < q \ll N = nM$ and avoid any $O(N)$ work.

1. Variational form, symmetry, and positive definiteness. Let $P \equiv SS^\top \in \mathbb{R}^{N \times N}$ be the diagonal "mask" matrix that keeps observed entries and zeros missing ones (so $P = P^\top = P^2$). The system (1) is the normal equation for the regularized least-squares objective

$$\min_{W \in \mathbb{R}^{n \times r}} \frac{1}{2} \left\| S^\top \mathrm{vec}(T) - S^\top \mathrm{vec}(KWZ^\top) \right\|_2^2 + \frac{\lambda}{2} \mathrm{Tr}(W^\top KW) = \frac{1}{2} \| P \circ (T - KWZ^\top) \|_F^2 + \frac{\lambda}{2} \mathrm{Tr}(W^\top KW), \tag{2}$$

where $\circ$ denotes Hadamard product and we used $\mathrm{Tr}(W^\top KW) = \|W\|_{\mathcal{H}}^2$ as the RKHS penalty. Writing the first term as $\frac{1}{2} \| P^{1/2}(\mathrm{vec}(T) - (Z \otimes K)\,\mathrm{vec}(W)) \|_2^2$ shows the Hessian is

$$A \equiv (Z \otimes K)^\top P(Z \otimes K) + \lambda(I_r \otimes K) \in \mathbb{R}^{nr \times nr}, \qquad b \equiv (I_r \otimes K)\, \mathrm{vec}(B) = \mathrm{vec}(KB).$$

Thus A is symmetric. If $K \succ 0$ and $\lambda > 0$, then $A \succ 0$ because for any $x \neq 0$,

$$x^\top Ax = \| P^{1/2}(Z \otimes K)x \|_2^2 + \lambda\, x^\top (I_r \otimes K)x \geq \lambda\, x^\top (I_r \otimes K)x > 0.$$

(If K is only psd, one can add a nugget εI_n to K or instead regularize with $\lambda(I_r \otimes I_n)$; we proceed with the $K \succ 0$ case.) Hence we can solve (1) with (preconditioned) conjugate gradients (CG/PCG).

2. Why PCG helps. Direct solution would require forming A and performing a dense factorization costing $O((nr)^3) = O(n^3 r^3)$. In contrast, PCG requires only: (i) repeated matrix–vector products $y \leftarrow Ax$ and (ii) repeated applications of a preconditioner M^{-1}, with overall cost $\approx$ #iters $\times$ (matvec + precond). Our goal is to implement both in $O(n^2 r + qr)$ time per iteration and memory $O(nr + qr)$, never touching N-scale arrays.

3. PCG (brief). Choose an SPD preconditioner $M \approx A$ that is cheap to invert. Starting from x_0 (often 0 or the previous ALS iterate), define $r_0 = b - Ax_0$ and solve $Mz_0 = r_0$. Set $p_0 = z_0$ and for $t = 0, 1, 2, \ldots$ iterate

$$\alpha_t = \frac{\langle r_t, z_t \rangle}{\langle p_t, Ap_t \rangle}, \quad x_{t+1} = x_t + \alpha_t p_t, \quad r_{t+1} = r_t - \alpha_t Ap_t, \quad \text{solve } Mz_{t+1} = r_{t+1}, \quad \beta_t = \frac{\langle r_{t+1}, z_{t+1} \rangle}{\langle r_t, z_t \rangle}, \quad p_{t+1} = z_{t+1} + \beta_t$$

The algorithm only needs the ability to compute Ap_t (matvec) and to apply M^{-1}.

Efficient matrix–vector products without forming A

Represent an input vector $x \in \mathbb{R}^{nr}$ as a matrix $X \in \mathbb{R}^{n \times r}$ such that $x = \mathrm{vec}(X)$. Use the identity

$$(Z \otimes K)\, \mathrm{vec}(X) = \mathrm{vec}(KXZ^\top), \tag{3}$$

which is a special case of $\mathrm{vec}(AXB^\top) = (B \otimes A)\, \mathrm{vec}(X)$.

Observed-entry operator implemented with index lists

Let the observed entries of the mode-k unfolding be indexed by pairs (i_t, j_t) for $t = 1, \ldots, q$ with $i_t \in [n]$ and $j_t \in [M]$. Then for any matrix $U \in \mathbb{R}^{n \times M}$,

$$S^\top \text{vec}(U) = \left(U_{i_t, j_t} \right)_{t=1}^q \in \mathbb{R}^q, \qquad \text{and} \qquad \text{reshape}_{n \times M}(Sv) \text{ has nonzeros } (i_t, j_t) \text{ equal to } v_t.$$

Thus we can realize $S^\top$ (gather) and S (scatter) in $O(q)$ time using stored index arrays (i_t, j_t).

Avoiding explicit formation of Z. Although $Z \in \mathbb{R}^{M \times r}$ is defined as a Khatri–Rao product, $M = \prod_{i \neq k} n_i$ can be enormous, so we do *not* store Z. Instead, for each observed tensor entry we typically store its full multi-index $(i_1^{(t)}, \ldots, i_d^{(t)})$; the corresponding row needed in (4) and (7) is

$$Z_{j_t, :} = A_d(i_d^{(t)}, :) \odot \cdots \odot A_{k+1}(i_{k+1}^{(t)}, :) \odot A_{k-1}(i_{k-1}^{(t)}, :) \odot \cdots \odot A_1(i_1^{(t)}, :),$$

which can be computed on the fly in $O((d-1)r)$ time per observed entry (or faster if intermediate Hadamard products are cached). This keeps both memory and time independent of M.

Matvec formula

Given $X \in \mathbb{R}^{n \times r}$, compute $Y \in \mathbb{R}^{n \times r}$ so that $\text{vec}(Y) = A \text{vec}(X)$. Write $G \equiv KX \in \mathbb{R}^{n \times r}$. Then the observed predicted entries of $U \equiv KXZ^\top \in \mathbb{R}^{n \times M}$ are

$$u_t \equiv U_{i_t, j_t} = G_{i_t, :} \cdot Z_{j_t, :} \quad (t = 1, \ldots, q), \tag{4}$$

each a length-r dot product. Now form the sparse matrix $\widetilde{U} \in \mathbb{R}^{n \times M}$ with $(\widetilde{U})_{i_t, j_t} = u_t$ and all other entries zero (this is exactly $\text{reshape}(SS^\top \text{vec}(U))$). Finally apply $(Z \otimes K)^\top$ using the transpose identity

$$(Z \otimes K)^\top \text{vec}(\widetilde{U}) = \text{vec}(K \widetilde{U} Z), \tag{5}$$

obtaining the main term $K(\widetilde{U}Z)$. Adding the Tikhonov term gives

$$Y = K(\widetilde{U}Z) + \lambda KX. \tag{6}$$

How to compute $\widetilde{U}Z$ in $O(qr)$

We never form $\widetilde{U}$ explicitly as an $n \times M$ array. Instead, compute the product $H \equiv \widetilde{U}Z \in \mathbb{R}^{n \times r}$ by accumulating contributions from the q nonzeros:

$$H_{i_t, :} \mathrel{+}= u_t Z_{j_t, :} \qquad (t = 1, \ldots, q). \tag{7}$$

Each update is a SAXPY of length r, so the cost is $O(qr)$. Then compute KH in $O(n^2 r)$ time (dense K), and add λKX.

Implementation sketch (matvec). Given $X \in \mathbb{R}^{n \times r}$, precompute $G \leftarrow KX$. Initialize $H \leftarrow 0 \in \mathbb{R}^{n \times r}$. Loop over observed entries $t = 1, \ldots, q$: compute (or fetch) the Khatri–Rao row $z_t \equiv Z_{j_t, :}$, compute the scalar $u_t \leftarrow G_{i_t, :} z_t^\top$, and update $H_{i_t, :} \leftarrow H_{i_t, :} + u_t z_t$. Finally return $Y \leftarrow KH + \lambda G$.

Matvec complexity.   Assuming dense K:

- $G = KX$: $O(n^2 r)$.

- gather u_t via (4): $O(qr)$ given access to $Z_{j_t,:}$; if $Z_{j_t,:}$ is computed on the fly from the CP factors, add $O(q(d-1)r)$.

- accumulate $H = \widetilde{U}Z$ via (7): $O(qr)$ (plus the same cost to form $Z_{j_t,:}$ if needed).

- KH and λKX: $O(n^2 r)$ (can reuse G for KX).

Total per matvec: $O(n^2 r + qr)$ time if Z-rows are available (or $O(n^2 r + qdr)$ if computed on the fly), $O(nr + qr)$ memory for X, G, H and index lists; crucially independent of $N = nM$.

Computing the right-hand side b without forming T

The right-hand side is $b = \mathrm{vec}(KB)$. We can compute $B = TZ$ using only observed entries: if the observed tensor value at (i_t, j_t) is t_t, then

$$B_{i_t,:} \mathrel{+}= t_t\, Z_{j_t,:},$$

which costs $O(qr)$ given access to $Z_{j_t,:}$ (or $O(qdr)$ if each $Z_{j_t,:}$ is formed on the fly from the CP factors), followed by KB in $O(n^2 r)$. (This is the same sparse accumulation pattern as (7).)

Preconditioning

Kronecker "full-observation" preconditioner

A standard and effective choice is to drop the mask P (equivalently, pretend all entries are observed). Then

$$A_0 \equiv (Z \otimes K)^\top (Z \otimes K) + \lambda(I_r \otimes K) = (Z^\top Z) \otimes (K^2) + \lambda(I_r \otimes K).$$

Let $G \equiv Z^\top Z \in \mathbb{R}^{r \times r}$. Since Z is a Khatri–Rao product, G can be computed without forming Z via the Hadamard product (denoted $*$) of Gram matrices,

$$G = \underset{i \neq k}{*} (A_i^\top A_i),$$

costing $O(\sum_{i \neq k} n_i r^2)$. If we precompute eigendecompositions

$$K = U \Lambda U^\top, \qquad G = V \Sigma V^\top,$$

then A_0 diagonalizes in the Kronecker basis:

$$A_0 = (V \otimes U) \, \mathrm{diag}\big(\sigma_a \lambda_b^2 + \lambda \lambda_b\big)_{a \in [r], b \in [n]} (V \otimes U)^\top.$$

Hence applying $M^{-1} \approx A_0^{-1}$ to a vector $x = \mathrm{vec}(X)$ can be done by:

1. transform $\widehat{X} \leftarrow U^\top X V$ (two small dense multiplies),

2. elementwise divide $\widehat{X}_{b,a} \leftarrow \widehat{X}_{b,a} / (\sigma_a \lambda_b^2 + \lambda \lambda_b)$,

3. inverse transform $X \leftarrow U \widehat{X} V^\top$.

This costs $O(n^2 r + n r^2)$ per application (often dominated by $O(n^2 r)$ when $n \geq r$), after one-time setup $O(n^3 + r^3)$.

Why this is reasonable. A equals A_0 when $P = I_N$; when entries are missing uniformly at random, $\mathbb{E}[P] = (q/N)I_N$, so $\mathbb{E}[(Z \otimes K)^\top P(Z \otimes K)] = (q/N)(Z^\top Z \otimes K^2)$ is a scaled version of the same Kronecker term. Thus A_0 (or the scaled variant αA_0 with $\alpha = q/N$) captures the dominant spectral structure, while P acts as a perturbation; PCG corrects the mismatch through iterations.

Simpler (cheaper) preconditioners

If eigendecompositions are too costly, a cheaper alternative is a block-diagonal preconditioner

$$M_{\mathrm{bd}} = (\mathrm{diag}(G) \otimes K^2) + \lambda(I_r \otimes K),$$

which decouples the r components, requiring r solves with $n \times n$ matrices of the form $(g_{\ell\ell}K^2 + \lambda K) = K(g_{\ell\ell}K + \lambda I)$. If K is factored once (Cholesky), these are fast; if K is large, one can use an approximate factorization (pivoted Cholesky / incomplete Cholesky) as a preconditioner for these inner solves.

Overall complexity and scaling

Let m be the number of PCG iterations to reach a desired tolerance; standard theory gives $m = O\big(\sqrt{\kappa(M^{-1}A)} \log(1/\varepsilon)\big)$ for relative error ε, so a good preconditioner aims to make $\kappa(M^{-1}A)$ close to 1.

- One-time setup: compute Z and (optionally) $G = Z^\top Z$ in $O(Mr^2)$ if done naively, but in CP-ALS contexts Z is implicit and G is usually assembled from Hadamard products of Gram matrices of each factor at cost $O(\sum_{i \neq k} n_i r^2)$, avoiding M.

- Right-hand side: compute B in $O(qr)$ and then KB in $O(n^2r)$.

- Each PCG iteration:

  - matvec Ax: $O(n^2r + qr)$.
  - preconditioner apply (Kronecker-eig): $O(n^2r + nr^2)$.
  - vector updates/inner products: $O(nr)$.

Hence total time $O\big(m(n^2r + qr + nr^2) + qr + n^2r\big)$, which is dramatically better than $O(n^3r^3)$ when $m \ll n^2r^2$ and $q \ll N$.

Key point: no N-scale work. All operations are expressed in terms of n, r, q and small Gram matrices; selection/scatter uses only the q observed indices and values.